Plesae check the GitHub repo for better readabilty

https://github.com/ibrahimmenshawy94/PG-DevOps-Project---Hotel-Side-Hospital.git

Introduction:

Hotel-Side Hospital, a globally renowned hospital chain headquartered in Australia, is aiming to streamline its operation by setting up an infrastructure within the hotel premises. However, in order to maintain seamless functioning and scalability, they require fully managed virtual machines (VMs) on the Amazon Web Services (AWS) platform. The organization seeks an automated provisioned infrastructure solution that can enable them to effortlessly create new Amazon Elastic Kubernetes Service (EKS) clusters, whenever required, and promptly delete them when they are no longer needed. This will optimize resource allocation and enhance operational efficiency.

Prerequisites:

Skills used in the project and their usage in the industry are as below:

• Terraform: It is an infrastructure-as-code tool that allows you to define and provision resources in a cloud environment. In this project, Terraform is used to define the infrastructure components and manage their lifecycle.

• EKS: It is a managed Kubernetes service provided by AWS. In this project, an EKS cluster is created using Terraform, which provides a scalable and highly available environment for running containerized workloads.

• EC2: EC2 instances are virtual servers provided by AWS. In this project, EC2 instances are provisioned using Terraform, which allows you to specify the desired instance types, operating systems, and other configurations.

Repository Setup:

- Steps to initialize the Git repository.

mkdir HSH

git init
git remote add origin https://github.com/ibrahimmenshawy94/PG-DevOps-Project---Hotel-Side-Hospital.git

git pull origin main

```
root@ip-172-31-68-179:~# mkdir HSH
root@ip-172-31-68-179:~# cd HSH
root@ip-172-31-68-179:~/HSH# git init
hint: Using 'master' as the name for the initial branch. This default branch name
hint: is subject to change. To configure the initial branch name to use in all
hint: of your new repositories, which will suppress this warning, call:
hint:
hint:   git config --global init.defaultBranch <name>
hint:
hint: Names commonly chosen instead of 'master' are 'main', 'trunk' and
hint: 'development'. The just-created branch can be renamed via this command:
hint:
hint:   git branch -m <name>
Initialized empty Git repository in /root/HSH/.git/
root@ip-172-31-68-179:~/HSH# git remote add origin https://github.com/ibrahimmenshawy94/PG-DevOps-Project---Hotel-Side-Hospital.git
root@ip-172-31-68-179:~/HSH#
```

Terraform Configuration:

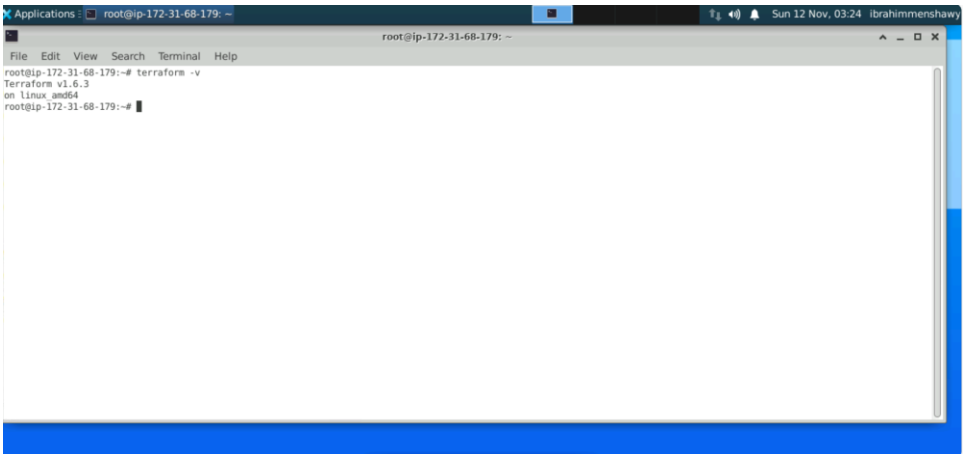

AWS CLI

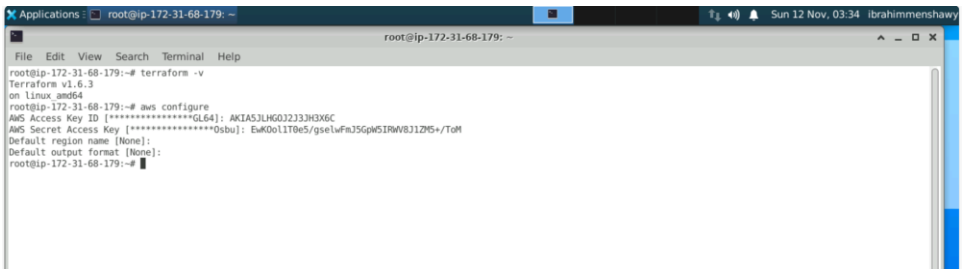

Terraform files:

1- **provider.tf**

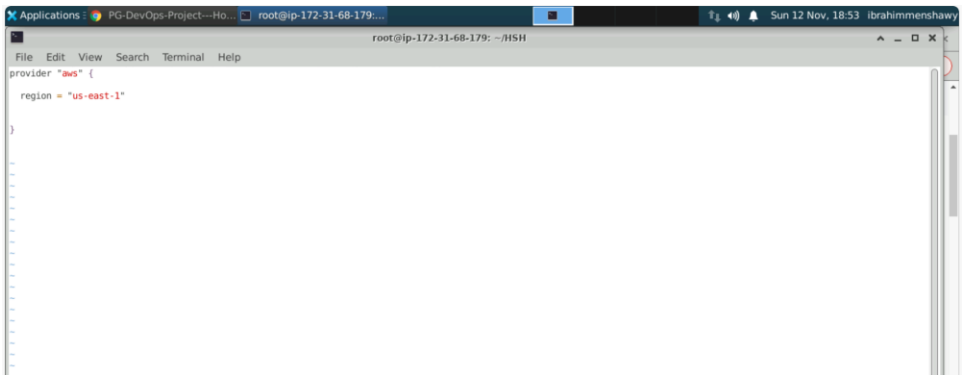


```
provider "aws" {

  region = "us-east-1"


}
```

2 -

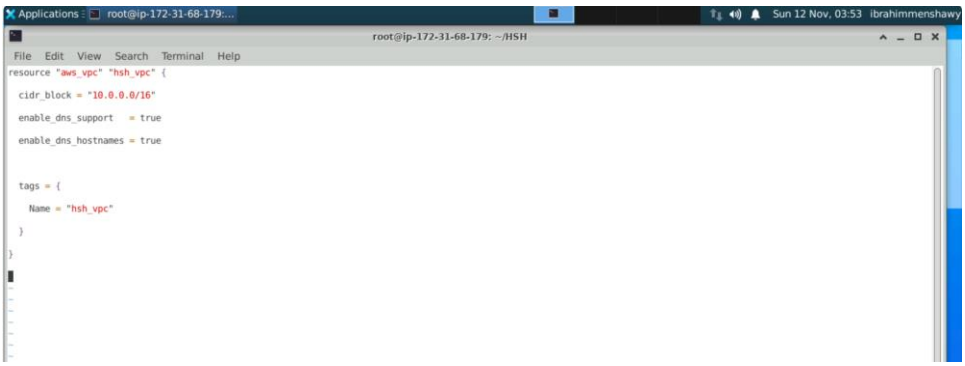


```
resource "aws_vpc" "hsh_vpc" {

  cidr_block = "10.0.0.0/16"

  enable_dns_support   = true
```

enable_dns_hostnames = true

tags = {

Name = "hsh_vpc"

}

}

3 - subnets.tf

esource "aws_subnet" "hsh_public_subnet" {

vpc_id = aws_vpc.hsh_vpc.id

```
  cidr_block        = "10.0.1.0/24"

  map_public_ip_on_launch = true




  tags = {

    Name = "hsh_public_subnet"

  }

}
```

4 – routetable.tf

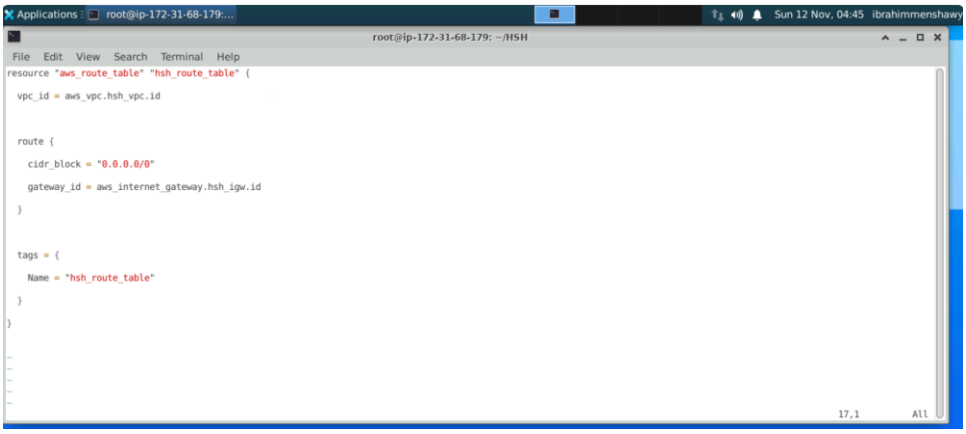


```
resource "aws_route_table" "hsh_route_table" {

  vpc_id = aws_vpc.hsh_vpc.id
```

```
 route {

   cidr_block = "0.0.0.0/0"

   gateway_id = aws_internet_gateway.hsh_igw.id

 }




 tags = {

   Name = "hsh_route_table"

 }

}
```

5 - Route_table_association_with_public_subnets.tf

```
resource "aws_route_table_association" "hsh_rta_public_subnet" {

  subnet_id      = aws_subnet.hsh_public_subnet.id

  route_table_id = aws_route_table.hsh_route_table.id

}
```

6 - internetgateway.tf

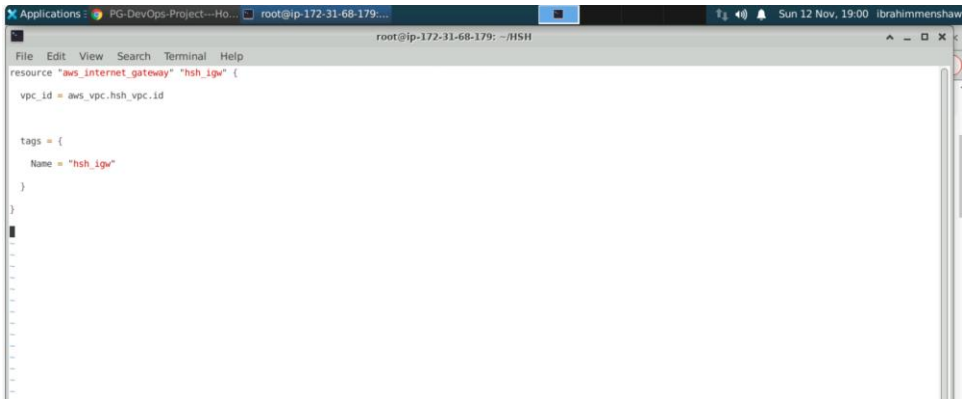


```
resource "aws_internet_gateway" "hsh_igw" {

  vpc_id = aws_vpc.hsh_vpc.id



  tags = {

    Name = "hsh_igw"

  }
```

}

7 - autoscaling.tf

resource "aws_launch_configuration" "hsh_lc" {

name_prefix = "hsh-lc-"

image_id = "ami-0fc5d935ebf8bc3bc"

instance_type = "t2.micro"

security_groups = [aws_security_group.hsh_sg.id]

lifecycle {

create_before_destroy = true

```
  }

}



resource "aws_autoscaling_group" "hsh_asg" {

  launch_configuration = aws_launch_configuration.hsh_lc.id

  vpc_zone_identifier  = [aws_subnet.hsh_public_subnet.id]

  max_size             = 3

  min_size             = 1

  desired_capacity     = 1



  tag {

    key                 = "Name"

    value               = "hsh-instance"

    propagate_at_launch = true
```

```
  }

}



resource "aws_autoscaling_policy" "hsh_target_tracking" {

  name                   = "hsh-cpu-target-tracking"

  autoscaling_group_name = aws_autoscaling_group.hsh_asg.name

  policy_type            = "TargetTrackingScaling"

  estimated_instance_warmup = 200



  target_tracking_configuration {

    predefined_metric_specification {

      predefined_metric_type = "ASGAverageCPUUtilization"

    }

    target_value = 30.0

  }
```

}

git add .

git commit -m "Initial commit or a description of changes"

git push -u origin main

Jenkins congiuration:

1- Add terraform plugin: myterraform

1- Create a credentials for AWS IAM

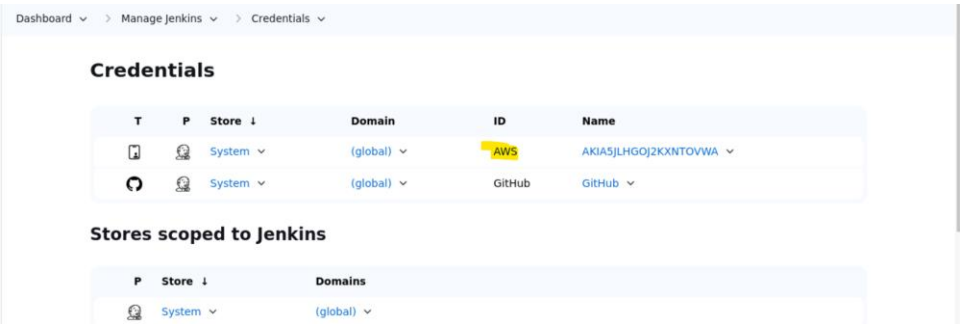

- Create a new pipeline that checkout the github repo and apply the terraform scripts:

```
pipeline {
  agent any
  tools {
    terraform 'myterraform'
  }

  stages {
    stage('Checkout') {
      steps {
        checkout([$class: 'GitSCM', branches: [[name: '*/main']], extensions: [], userRemoteConfigs: [[url: 'https://github.com/ibrahimmenshawy94/PG-DevOps-Project---Hotel-Side-Hospital.git']]])
      }
    }
    stage('Terraform init') {
      steps {
        sh 'terraform init'
      }
    }
    stage('Terraform apply') {
      steps {
        withCredentials([[
          $class: 'AmazonWebServicesCredentialsBinding',
          credentialsId: 'AWS'
        ]]) {
          sh '''
            export AWS_ACCESS_KEY_ID=$AWS_ACCESS_KEY_ID
            export AWS_SECRET_ACCESS_KEY=$AWS_SECRET_ACCESS_KEY
```

```
                export AWS_DEFAULT_REGION='us-east-1'
                terraform apply -auto-approve
            '''
        }
    }
  }
 }
}
```

- Build the terraform pipeline:

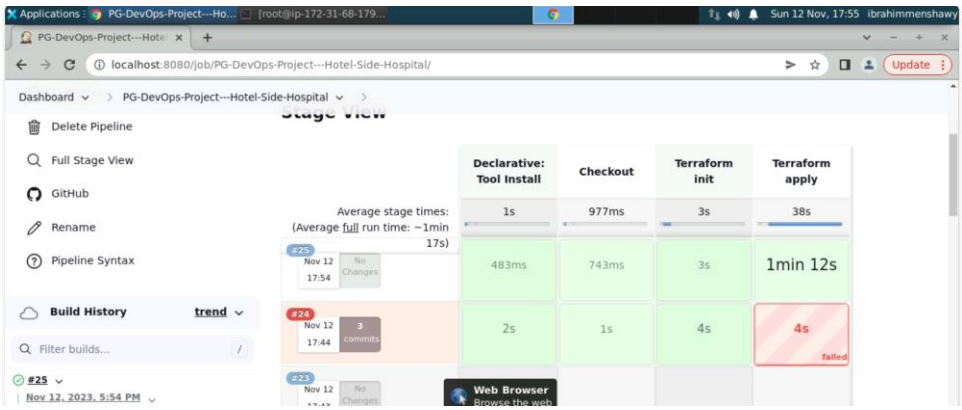

Check the instance on AWS

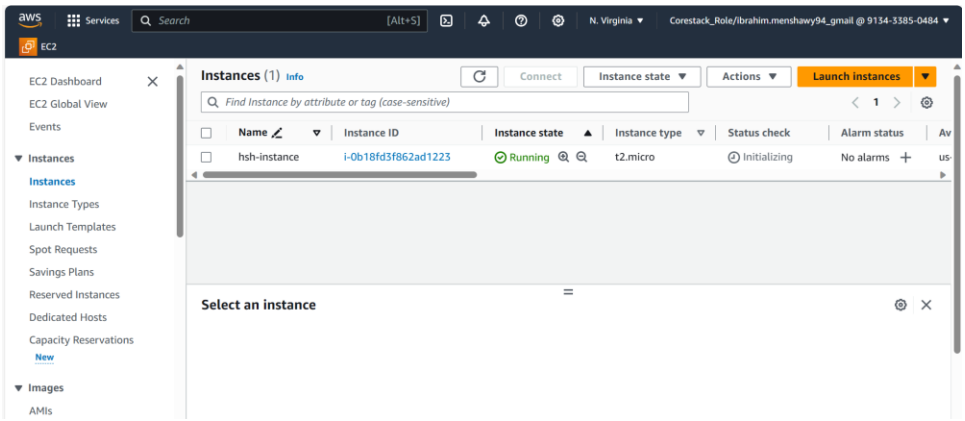

- Check the Autoscaling group:

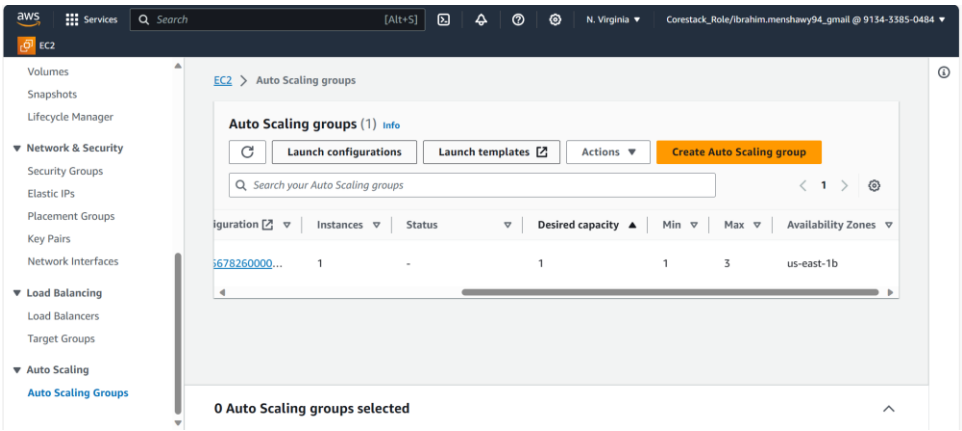

CPU usage before installing stress utility:

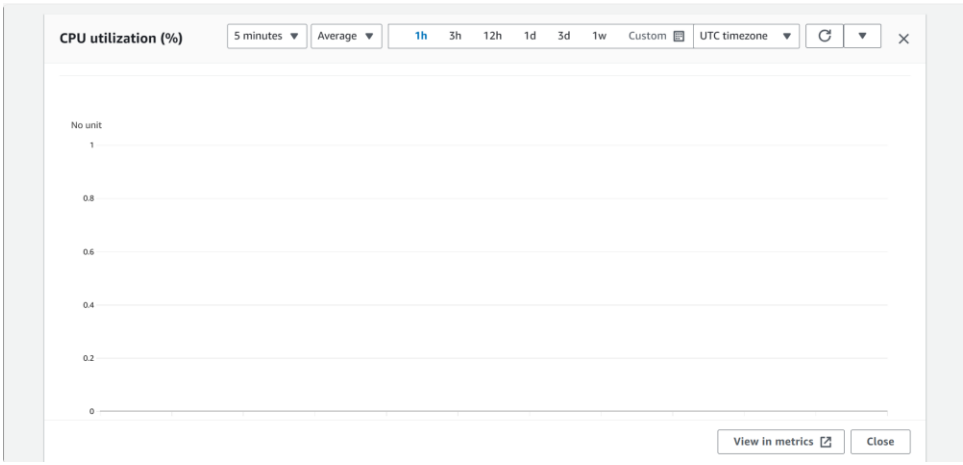

Installing stress utility on the instance (I used Ubuntu instance hence the difference in command)

sudo apt update

sudo apt install stress -y

sudo stress --cpu 8 --verbose --timeout 3000s

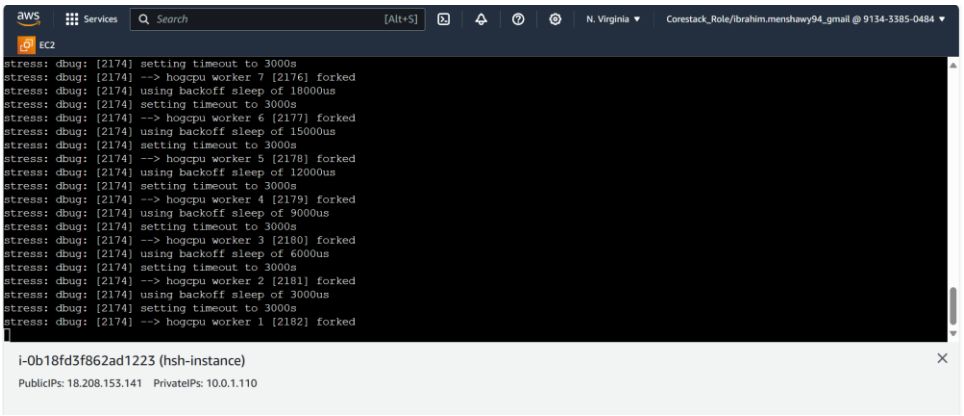

- Check the status of the instances afterr the triger is applied (cpu > 20%)

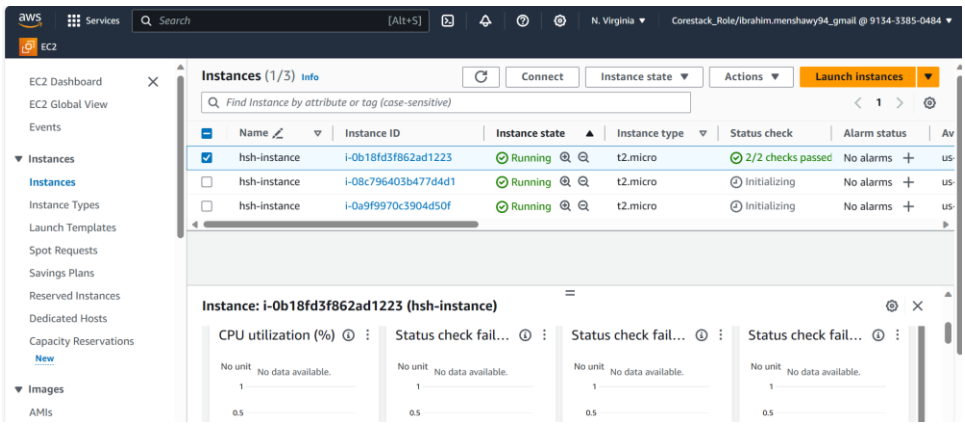

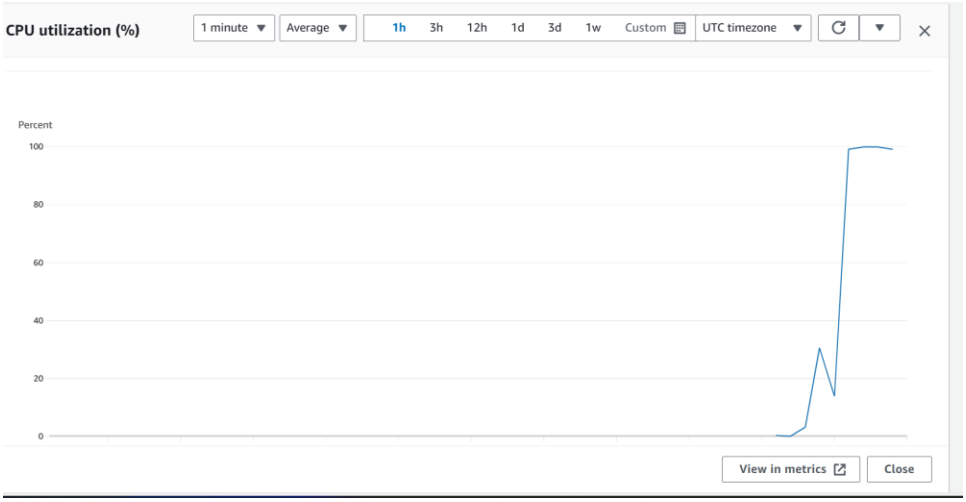
CPU utilization (%)
1 minute
Average
1h
3h
12h
1d
3d
1w
Custom
UTC timezone
Percent
100
80
60
40
20
0
View in metrics
Close

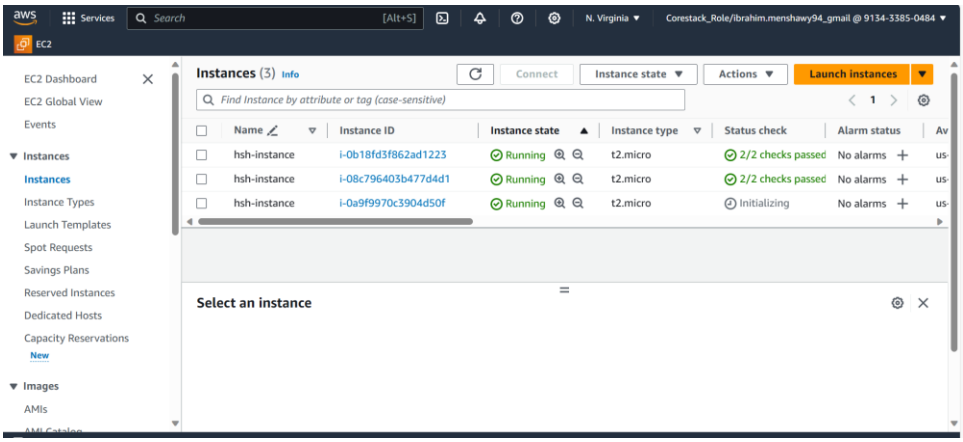
Services
Search
[Alt+S]
N. Virginia
Corestack_Role/ibrahim.menshawy94_gmail @ 9134-3385-0484
EC2
EC2 Dashboard
EC2 Global View
Events
Instances
Instances
Instance Types
Launch Templates
Spot Requests
Savings Plans
Reserved Instances
Dedicated Hosts
Capacity Reservations
New
Images
AMIs
Instances (3) Info
Connect
Instance state
Actions
Launch instances
Find Instance by attribute or tag (case-sensitive)
Name
Instance ID
Instance state
Instance type
Status check
Alarm status
hsh-instance
i-0b18fd3f862ad1223
Running
t2.micro
2/2 checks passed
No alarms
hsh-instance
i-08c796403b477d4d1
Running
t2.micro
2/2 checks passed
No alarms
hsh-instance
i-0a9f9970c3904d50f
Running
t2.micro
Initializing
No alarms
Select an instance

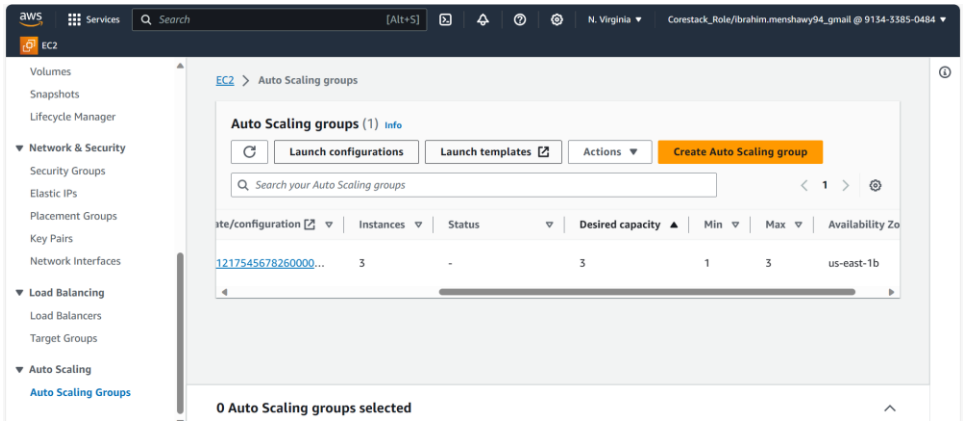
Services
Search
[Alt+S]
N. Virginia
Corestack_Role/ibrahim.menshawy94_gmail @ 9134-3385-0484
EC2
Volumes
Snapshots
Lifecycle Manager
Network & Security
Security Groups
Elastic IPs
Placement Groups
Key Pairs
Network Interfaces
Load Balancing
Load Balancers
Target Groups
Auto Scaling
Auto Scaling Groups
EC2 > Auto Scaling groups
Auto Scaling groups (1) Info
Launch configurations
Launch templates
Actions
Create Auto Scaling group
Search your Auto Scaling groups
Instances
Status
Desired capacity
Min
Max
Availability Zo
1217545678260000...
3
-
3
1
3
us-east-1b
0 Auto Scaling groups selected

- Now stop the stress utility

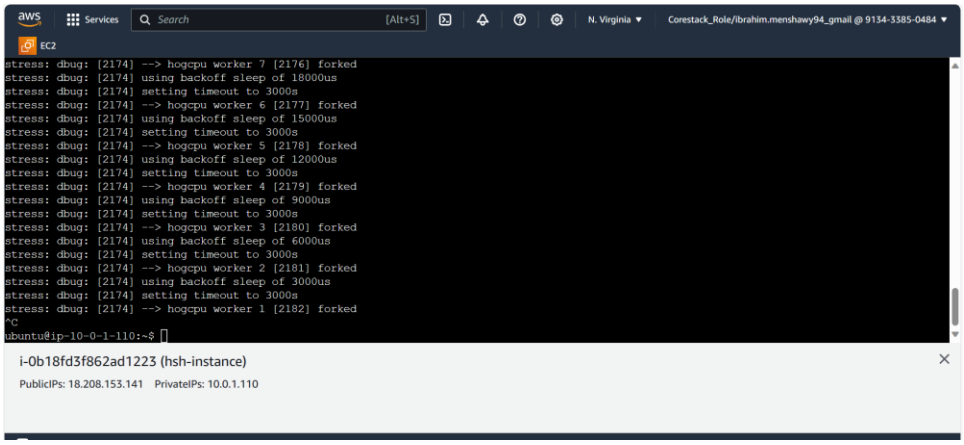

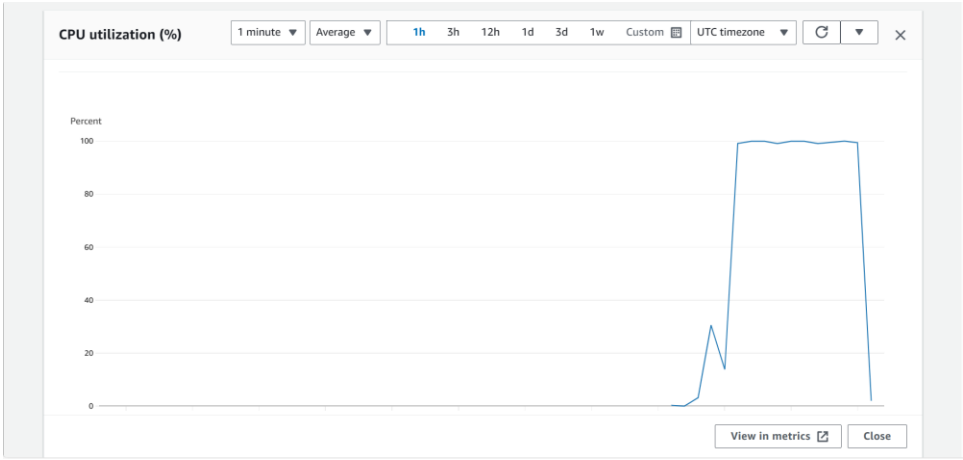

After some time, the group returnes to the desired state of 1 instance

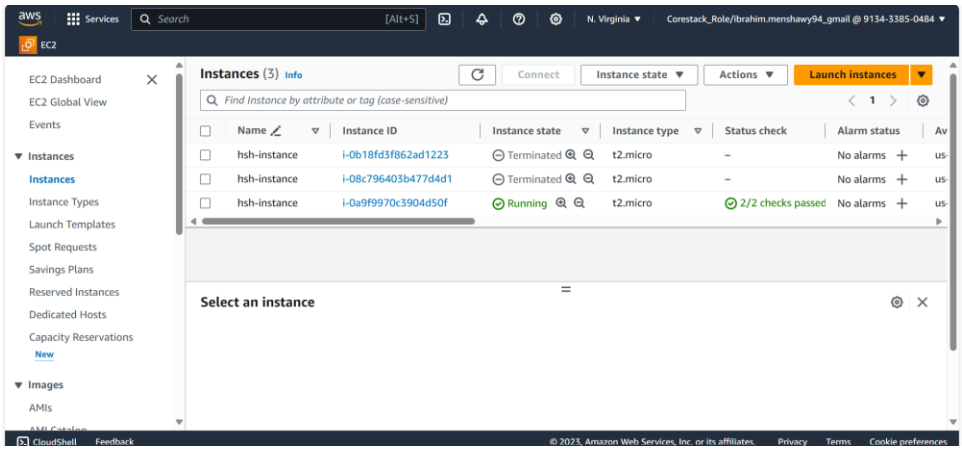

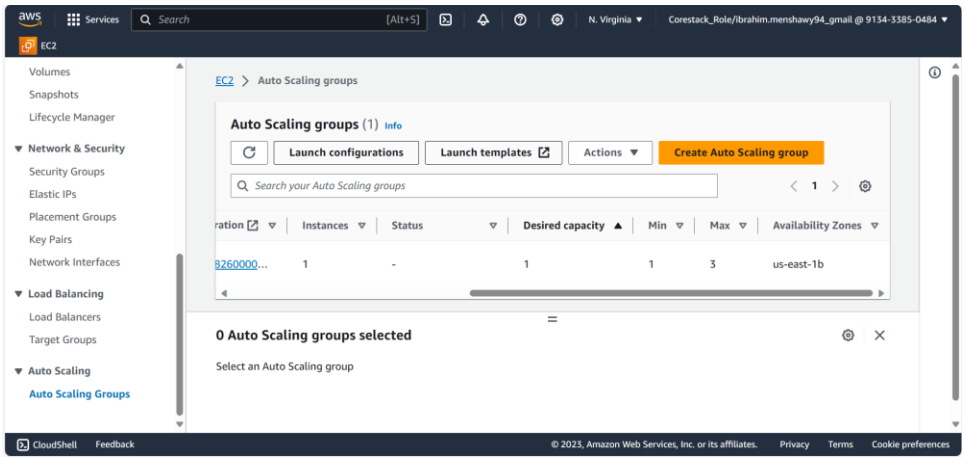

The END

Thank YOU